



The UK AI Skills Hub

A structural design gap — and what the evidence suggests is missing

Status: preliminary findings from an independent research project. The wider project remains at the scoping stage; an initial audience and pilot-unit direction have since been selected for testing (§6). Evidence and inference are distinguished throughout, and every source is tagged by its likely interest. This is not a finished proposal, nor a claim that the underlying problem has been exhaustively verified.

Overview/Editorial

The UK Government's AI Skills Hub is a national learning platform under Innovate UK and Skills England, delivered by PwC and industry partners. The scope of this initiative expanded in January 2026 to encompass all UK adults, with an ambition to upskill 10 million workers by 2030. This could almost pass as a proactive, uncharacteristically shrewd move from the government. Almost. There is a fundamental, independently corroborated flaw in the skills hub: it functions as a directory of pre-existing courses organised by technology vendor, not a structured programme mapped to the competencies employers and individuals actually need. If the government wants anything close to the proposed '1 in 7 people' in the country using this tool, they'll have to do better than vaguely pointing at a disjointed library of existing proprietary courses. An LSE academic analysis and two commercial reviewers tested the platform separately, converging on the same gaps: no competency mapping, no diagnostic-driven sequencing, no embedded practice, no support for organisational readiness.

The direction suggested here is grounded in existing research rather than invented from scratch. The same government programme that diagnosed the UK's AI skills gap has now published PRIMES, a framework describing evidence-informed training principles. Comparing the AI Skills hub platform against the PRIMES framework, reasonable conclusions can be made that the platform is at best ineffective, and at times directly contradictory to guidance. [institutional disconnect/lack of understanding argument here or in later section?]

1. The platform and the problem

The AI Skills Hub is a national training-navigation platform delivered under the BridgeAI programme, developed by Innovate UK in partnership with Skills England. It initially focused on four sectors. On 28 January 2026 it expanded to all UK adults, with a stated ambition to upskill 10 million workers by 2030 and a reported £4.1 million investment.

The core critique is structural, not about content quality: the platform curates existing third-party courses, but it doesn't sequence them against a competency framework, personalise pathways, or provide embedded practice. Reviewers call it "a directory, not a programme."



2. Independent corroboration

Three separately-authored reviews, none citing the others, converge on overlapping specific findings:

“A mixture of existing training content loosely connected by the topic of AI.”

— LSE Impact of Social Sciences blog (Cardoso-Silva & Brown), on the AI Skills Hub's course organisation

Source	Evidential role	Key finding
LSE Impact of Social Sciences blog (Cardoso-Silva & Brown) Independent academic — LSE Data Science Institute / Digital Skills Lab	Independent policy analysis	Courses organised by technology vendor, not by the competencies the government's own research identifies as most needed (72%+ of surveyed organisations: non-technical skills).
Techosaurus / Quantum Rise Commercial AI-training providers — competing interest (flagged)	Observed platform behaviour	Five flaws: no organisational-readiness support, vendor/platform lock-in, no practice sandbox, no diagnostics or scaffolding, disconnect from how people actually use AI.
The Human Co. Commercial AI-training/consulting — competing interest (flagged)	Observed platform behaviour	Rated the £4.1m platform 3/10 on design despite adequate course content; intermediate learners shown ~71% beginner-level material; no sequencing; progress tracking “nearly invisible.”

Note on interest: two of the three reviewers run commercial AI-training businesses that compete with the platform they reviewed, and are flagged accordingly. Their specific, hands-on structural findings nonetheless align closely with the independent academic critique — the corroboration is across source types, not just repetition of one view. That corroboration is about the broad structural problem, not every number. Figures like “71% beginner-level material” or the “3/10” rating come from individual reviews, not a replicated evaluation — read them as illustrative, not as measured consensus.

3. The underlying cause: “a directory, not a programme”

The pattern across all three reviews is the same: a library of courses, organised by technology or vendor rather than by skill. There's no diagnostic step to route learners to their level, no mandated sequencing, and no built-in low-stakes practice. One reviewer found intermediate learners shown roughly 71% beginner-level content — nearly doubling realistic completion time.

The irony is that this isn't a case of missing evidence. The same government-commissioned research programme (Dr Nisreen Ameen, Royal Holloway, University of London, with Skills England — 23 workshops, 10 employer case studies, a 536-response survey) has already published a framework for what effective AI training looks like, summarised below. None of the three reviews found evidence of competency mapping, diagnostic sequencing, or PRIMES alignment on the platform. That's a reasonable basis to infer the framework wasn't applied — though none of the reviewers saw the Hub's internal design documents, so it stays an inference, not a confirmed fact.

4. What good looks like: two reference models

4.1 PRIMES — the government's own evidence-based framework



Published 10 June 2026 as part of the same “What Works for AI Upskilling in the UK” research programme, PRIMES sets out six accreditation-style principles for effective AI training. Each row below adds one concrete, checkable requirement from the framework’s own criteria document, not just its six names:

Practical	Links to real tasks and decisions; recognises existing informal AI use rather than assuming a zero starting point; makes explicit when AI should and should not be used.
Reachable	Accessible in time, format and language; explicitly designed around intersecting barriers (income, age, disability, gender), not a single “average learner.”
Integrated	Fits into existing systems and standards; mandatory before staff use AI on regulated, confidential or safety-critical work.
Modular	Short, stackable units (30–90 minutes); learners can enter at different levels and return as needs change.
Expandable	Builds transferable skills usable across tools and roles, not single-vendor competence.
Sustainable	Reviewed on a set cadence (3- and 6-month revisits); tracks confidence and quality of judgement over time, not just enrolment.

Common pitfalls

The same guide names 13 common training pitfalls found across the organisations surveyed. Three are named specifically:

- Training is too fast, too technical, or assumes prior knowledge.
- Existing skills and informal AI use are not recognised.
- Training focuses on tools rather than transferable skills.

Where employers actually are

The same guide’s companion “AI Skills Adoption Pathway” shows survey respondents’ organisations still overwhelmingly concentrated in early stages:

Stage	Share of organisations
Awareness	21%
Exploration	19%
Integration	7%
Strategy	4%
Scaling	1%

Most respondents’ organisations remain concentrated in the earliest two stages (40% combined); only 12% have reached Integration, Strategy or Scaling combined.



4.2 Elements of AI — evidence a coherent alternative can scale in reach

Elements of AI, a free course from MinnaLearn and the University of Helsinki, offers a working counter-example: vendor-neutral, requiring no programming or technical prerequisite, structured as six short chapters combining theory, practical exercises and self-reflection, completed at the learner's own pace. It has reportedly reached approximately one million learners globally and is cited by the LSE reviewers as the model the UK's own platform should have followed. That's reach at scale — not proof of learning outcomes. Enrolment and completion are one thing; measured capability transfer into the workplace is a harder, different claim this doesn't establish.

5. Emerging design principles for a credible alternative

PRIMES, the platform reviews, and the Gradual Release of Responsibility model together suggest seven provisional design principles — starting points for Grounded AI Practice to test, not settled conclusions:

- 1. Modular** — Units of roughly 30–90 minutes with clear entry and re-entry points, rather than long, fixed-length courses.
- 2. Diagnostic-driven** — Sequencing that shifts responsibility from instructor to learner progressively — modelling, then guided practice, then independent practice — rather than jumping straight to unsupported application.
- 3. Practice-led** — Embedded, low-stakes practice (“sandbox” activity), not instructional content alone.
- 4. Bounded** — Explicit guidance on when AI should and should not be used, not only how to use it.
- 5. Transferable** — Tool-agnostic skills prioritised over single-vendor competence.
- 6. Revisitable** — Planned revisit points and outcome tracking (for example, 3- and 6-month check-ins) rather than one-off completion metrics.
- 7. Confidence-calibrating** — Training designed to build calibrated confidence, not just knowledge — see below.

6. Where this sits in the project

Grounded AI Practice's own direction has moved from open scoping to a set of working decisions since this research began:

Initial audience	Individuals seeking practical, everyday AI literacy, with particular attention to employees at small organisations who lack the employer-provided L&D infrastructure that PRIMES and the AI Skills Hub both assume.
First output	A single pilot learning unit — not a roadmap, course library or platform — sized to PRIMES' modular criterion (30–90 minutes) and sequenced using the Gradual Release of Responsibility pattern (§5), tested with real learners before any wider structure is built.
Pilot capability	Effective prompting, working title “Effective prompting — what's really happening when you hit send” — chosen to show a learner the gap between what they type and what the model actually does with it, while teaching an immediately usable skill.

None of this claims to have solved the problem above. A single pilot unit tests an approach; it doesn't prove one works. But each decision was made to dodge the exact failure mode this briefing documents — one small, well-sequenced unit is a more direct test than a library of loosely-connected content ever could be.

7. A note on method

Findings here draw on primary sources read directly wherever possible. Every source is tagged by its likely interest — government/official, independent/academic, vendor/commercial, or advocacy/membership body — so the evidence base's concentration is visible, not hidden. Where a finding would materially affect the project's direction, confirming and disconfirming evidence was sought together, not just evidence for the initial hypothesis. A few leads from this research (further critical commentary on the AI Skills Hub in the trade press, for example) haven't been independently verified yet, so they're left out as findings.

This briefing was first drafted 24 July 2026 and updated since with three further findings (PRIMES' full criteria, a named sequencing model, and confirm/disconfirm evidence on training and overconfidence) logged as the project's research continued — see RESEARCH_LOG.md Entries 026–028.

Specific limitations of this review, stated directly:

- No access to the Hub's internal design documentation — findings rest on external behaviour and published material, not internal design intent.
- No learner interviews or usage analytics beyond what the three reviews themselves report.
- No independent replication of the commercial reviewers' hands-on tests.
- Limited evidence of measured learning outcomes for either the Hub or the comparator models cited.
- The platform's design may continue to change after this briefing was written.
- This review evaluates the Hub's structure, not the quality of every individual course it curates.



Sources

Innovate UK / Skills England, UK AI Skills Hub (BridgeAI programme), expanded to all UK adults 28 January 2026. aiskillshub.org.uk

Cardoso-Silva, J. & Brown, J. (LSE), “The UK’s AI training ambition deserves better than a course directory,” LSE Impact of Social Sciences blog, 17 February 2026.

Quilter, S. (Techosaurus Ltd) & Farmer, E. (Quantum Rise), “We Tested the Government’s AI Skills Hub: Five Critical Flaws You Need to Know,” Techosaurus News, 30 January 2026.

Thomas, P. (The Human Co.), “The UK Government’s AI Skills Hub: A £4.1m Lesson in How Not to Build Real AI Capability,” thehumanco.org, 2026.

Ameen, N. (Royal Holloway, University of London / British Academy Policy-Led Innovation Fellowship, in partnership with Skills England), “What Works for AI Upskilling in the UK: Employer Guide,” 10 June 2026. gov.uk

Ameen, N. et al., “AI Skills for the UK Workforce,” Skills England / DfE, 30 October 2025 (companion AI Skills Tools Package, updated 4 November 2025). gov.uk

MinnaLearn & University of Helsinki, Elements of AI. elementsofai.com

Kentucky Department of Education, “Evidence-Based Instructional Practices: Explicit Teaching and Modeling” (Gradual Release of Responsibility framework). education.ky.gov

Kamali, Gerstner, Hullman & Groh, “Generative AI Literacy Training Improves Intelligence Analysts’ Discrimination of Real and AI-Generated Images,” arxiv.org/pdf/2606.28510

“Automation Bias in Large Language Model Assisted Diagnostic Reasoning Among AI-Trained Physicians,” medrxiv.org/content/10.1101/2025.08.23.25334280 (preprint, not peer-reviewed; primary source could not be independently verified).

Full source list, interest tagging, and detailed entry-by-entry findings are maintained in RESEARCH_LOG.md within the Grounded AI Practice project repository.

